

HOMWORK 4

Smit Shah

ID: 908 587 6440

<https://github.com/smit-1999/NaiveBayes>

Instructions: Use this latex file as a template to develop your homework. Submit your homework on time as a single pdf file to Canvas. Late submissions may not be accepted. Please wrap your code and upload to a public GitHub repo, then attach the link below the instructions so that we can access it. You can choose any programming language (i.e. python, R, or MATLAB). Please check Piazza for updates about the homework.

1 Best Prediction Under 0-1 Loss (10 pts)

Suppose the world generates a single observation $x \sim \text{multinomial}(\theta)$, where the parameter vector $\theta = (\theta_1, \dots, \theta_k)$ with $\theta_i \geq 0$ and $\sum_{i=1}^k \theta_i = 1$. Note $x \in \{1, \dots, k\}$. You know θ and want to predict x . Call your prediction \hat{x} . What is your expected 0-1 loss:

$$\mathbb{E}[\mathbb{1}_{\hat{x} \neq x}]$$

using the following two prediction strategies respectively? Prove your answer.

Strategy 1: $\hat{x} \in \arg \max_x \theta_x$, the outcome with the highest probability.

In this strategy, the prediction \hat{x} is the outcome with the highest probability according to the parameter vector i.e.

$\hat{x} \in \arg \max_x \theta_x$. Formula for expectation is given by $\sum_{x=1}^k x * p(x)$

\therefore Expected 0-1 loss will be:

$$\mathbb{E}[\mathbb{1}_{\hat{x} \neq x}]$$

=

$$E[1(\hat{x} \neq x)] = \sum_{i=1}^k P(\hat{x} \neq x | x = i) \cdot P(x = i)$$

Since \hat{x} is chosen as the outcome with the highest probability, $P(\hat{x} \neq x | x = i) = 1 - \theta_{max}$.

Therefore,

$$\begin{aligned} E[1(\hat{x} \neq x)] &= \sum_{i=1}^k (1 - \theta_{max}) \cdot \theta_{max} = \sum_{i=1}^k \theta_{max} - \sum_{i=1}^k \theta_{max}^2 \\ &= \theta_{max} * (1 - \theta_{max}) \end{aligned}$$

Strategy 2: You mimic the world by generating a prediction $\hat{x} \sim \text{multinomial}(\theta)$. (Hint: your randomness and the world's randomness are independent)

In this strategy, the prediction \hat{x} is generated from the multinomial distribution with the parameter vector θ . The randomness in this prediction and the randomness in the world's observation (x) are independent.

The expected 0-1 loss for this strategy is the probability that the predicted value \hat{x} is not equal to the true value x :

$$E[1(\hat{x} \neq x)] = P(\hat{x} \neq x) = 1 - P(\hat{x} = x)$$

In a multinomial distribution, the probability mass function is given by:

$$P(\hat{x} = i) = \theta_i$$

Therefore,

$$E[1(\hat{x} \neq x)] = 1 - P(\hat{x} = x) = 1 - \theta_x$$

So, the expected 0-1 loss for Strategy 2 is $1 - \theta_x$.

2 Best Prediction Under Different Misclassification Losses (6 pts)

Like in the previous question, the world generates a single observation $x \sim \text{multinomial}(\theta)$. Let $c_{ij} \geq 0$ denote the loss you incur, if $x = i$ but you predict $\hat{x} = j$, for $i, j \in \{1, \dots, k\}$. $c_{ii} = 0$ for all i . This is a way to generalize different costs on false positives vs false negatives from binary classification to multi-class classification. You want to minimize your expected loss:

$$\mathbb{E}[c_{x\hat{x}}]$$

Derive your optimal prediction \hat{x} .

To derive the optimal prediction \hat{x} that minimizes the expected loss $E[c_{x\hat{x}}]$, we need to find the value of \hat{x} that minimizes the expected loss over all possible values of \hat{x} . Mathematically, we want to minimize the following expression:

$$E[c_{x\hat{x}}] = \sum_{i=1}^k \sum_{j=1}^k c_{ij} \cdot P(x = i, \hat{x} = j)$$

Given that x follows a multinomial distribution with parameter vector θ , the joint probability $P(x = i, \hat{x} = j)$ can be expressed as $P(x = i) \cdot P(\hat{x} = j | x = i)$. Since \hat{x} is your prediction, we are interested in minimizing $E[c_{x\hat{x}}]$ with respect to \hat{x} . Therefore, the problem can be formulated as follows:

$$E[c_{x\hat{x}}] = \sum_{i=1}^k \sum_{j=1}^k c_{ij} \cdot P(x = i) \cdot P(\hat{x} = j | x = i)$$

To find the optimal \hat{x} , we minimize $E[c_{x\hat{x}}]$ with respect to \hat{x} by setting the derivative with respect to \hat{x} to zero. This gives us the following expression:

$$\frac{d}{d\hat{x}} E[c_{x\hat{x}}] = \sum_{i=1}^k c_{i\hat{x}} \cdot P(x = i) - \sum_{i=1}^k \sum_{j=1}^k c_{ij} \cdot P(x = i) \cdot P(\hat{x} = j | x = i) = 0$$

Solving this equation for \hat{x} gives the optimal prediction that minimizes the expected loss $E[c_{x\hat{x}}]$. Please note that the solution for \hat{x} might depend on the specific form of the cost matrix c_{ij} and the conditional probabilities $P(\hat{x} = j | x = i)$ which are derived from the multinomial distribution with parameter vector θ . The solution will vary based on the values in the cost matrix and the multinomial probabilities.

3 Language Identification with Naive Bayes (8 pts each)

Implement a character-based Naive Bayes classifier that classifies a document as English, Spanish, or Japanese - all written with the 26 lower case characters and space.

The dataset is languageID.tgz, unpack it. This dataset consists of 60 documents in English, Spanish and Japanese. The correct class label is the first character of the filename: $y \in \{e, j, s\}$. (Note: here each file is a document in corresponding language, and it is regarded as one data.)

We will be using a character-based multinomial Naïve Bayes model. You need to view each document as a bag of characters, including space. We have made sure that there are only 27 different types of printable characters (a to z, and space) – there may be additional control characters such as new-line, please ignore those. Your vocabulary will be these 27 character types. (Note: not word types!)

1. Use files 0.txt to 9.txt in each language as the training data. Estimate the prior probabilities $\hat{p}(y = e)$, $\hat{p}(y = j)$, $\hat{p}(y = s)$ using additive smoothing with parameter $\frac{1}{2}$. Give the formula for additive smoothing with parameter $\frac{1}{2}$ in this case. Print and include in final report the prior probabilities. (Hint: Store all probabilities here and below in $\log()$ internally to avoid underflow. This also means you need to do arithmetic in log-space. But answer questions with probability, not log probability.)

$$\hat{p}(y = e) =$$

$$(b_j + \alpha) / ((\sum_{j=1}^k b_j) + \alpha * k)$$

Here $\alpha = 0.5$ since it is given in question and $k = 3$ since y has 3 possible values. The same can be extrapolated for other priors. The answer we get for $\hat{p}(y = e) = 10.5/31.5 = 0.333$, and it will be same for other priors.

```

760HW4.pdf  nn.py 1, M  main.py M X
main.py > ...
39
40
41 def calculate_prior(df: pd.DataFrame):
42     # laplace smoothing factor
43     alpha = 0.5
44     total = len(train) + 3*alpha
45     e_prior = (
46         len(df.loc[df['file_name'].astype(str).str[0] == 'e']) + alpha) / total
47     s_prior = (
48         len(df.loc[df['file_name'].astype(str).str[0] == 's']) + alpha)/total
49     j_prior = (
50         len(df.loc[df['file_name'].astype(str).str[0] == 'j']) + alpha)/total
51
52     print('Prior probability values for e, j ,s with laplace factor are',
53         e_prior, j_prior, s_prior)
54     return math.log(e_prior), math.log(j_prior), math.log(s_prior)
55
56
57 def calculate_class_conditional(df: pd.DataFrame, label: string):
58     # laplace smoothing factor
59     alpha = 0.5
60
61     # in whole training dataframe filter out rows for the particular label

```

PROBLEMS 1 OUTPUT DEBUG CONSOLE TERMINAL PORTS

```

s 0 0 10
smitshah@Smit-Shah NaiveBayes % python3 main.py
Prior probability values for e, j ,s with laplace factor are 0.3333333333333333 0.3333333333333333 0.3333333333333333

```

2. Using the same training data, estimate the class conditional probability (multinomial parameter) for English

$$\theta_{i,e} := \hat{p}(c_i \mid y = e)$$

where c_i is the i -th character. That is, $c_1 = a, \dots, c_{26} = z, c_{27} = \text{space}$. Again use additive smoothing with parameter $\frac{1}{2}$. Give the formula for additive smoothing with parameter $\frac{1}{2}$ in this case. Print θ_e and include in final report which is a vector with 27 elements.

```

main.py > calculate_class_conditional
68 # print('Dataframe after filtering only letters and spaces as the features', df)
69
70 likelihood = {}
71 likelihood = dict.fromkeys(df.columns, 0)
72 total = df.to_numpy().sum() + (alpha*27)
73 print('Total sum of dataframe along with laplace smoothing for label',
74       label, 'is', total)
75
76 for col in df.columns:
77     col_sum = df[col].to_numpy().sum() + alpha
78     prob = col_sum/total
79     log_prob = math.log(prob)
80     likelihood[col] = (col_sum, prob, log_prob)
81 print('Likelihood dictionary for label', label, 'is', likelihood)
82
83 # sum of probabilities should be 1
84 res = tuple(sum(x) for x in zip(*likelihood.values()))
85 print('Sum of tuple values of likelihood dictionary for label', label, 'is', res)
86 return likelihood
87
88
89 def get_test_document_statistics(directory: string, filename: string, df: pd.DataFrame, conditional: pd.DataFrame):

```

PROBLEMS 1 OUTPUT DEBUG CONSOLE TERMINAL PORTS

```

Total sum of dataframe along with laplace smoothing for label e is 15132.5
Likelihood dictionary for label e is {'a': (910.5, 0.0601685114819098, -2.8106061285981796), 'b': (168.5, 0.011134974392863043, -4.497664277
973426), 'c': (325.5, 0.021509995043779945, -3.8392375661172182), 'd': (332.5, 0.021972575582355856, -3.8179601676699337), 'e': (1594.5, 0.1
053692383941847, -2.25028454131573), 'f': (286.5, 0.018932760614571286, -3.966861491611048), 'g': (264.5, 0.017478936064761277, -4.046758776
467488), 'h': (714.5, 0.047216256401784236, -3.05301703039592), 'i': (838.5, 0.055410540227986124, -2.892985446502579), 'j': (21.5, 0.001420
783082768875, -6.556547092632226), 'k': (56.5, 0.0037336857756484387, -5.590359389613447), 'l': (438.5, 0.028977366595076822, -3.54124021595
36045), 'm': (310.5, 0.020518751032545846, -3.8864161263923087), 'n': (876.5, 0.057921691723112505, -2.848663323404553), 'o': (975.5, 0.0644
6390219725756, -2.741649867702676), 'p': (253.5, 0.01675202378985627, -4.089236204734605), 'q': (8.5, 0.0005617049396993227, -7.484533864269
571), 'r': (814.5, 0.053824549810011564, -2.9220255997237423), 's': (1001.5, 0.06618205848339666, -2.7153458726599693), 't': (1212.5, 0.0801
2555757475633, -2.524160404954204), 'u': (403.5, 0.026664463902197257, -3.624423540055839), 'v': (140.5, 0.009284652238559392, -4.6793925389
92042), 'w': (234.5, 0.015496448042293078, -4.167144439879508), 'x': (17.5, 0.001156451346439782, -6.7623991468363736), 'y': (209.5, 0.01384
4374690236246, -4.27987628840365), 'z': (9.5, 0.0006277878737815959, -7.373308229159347), 'space': (2712.5, 0.1792499586981662, -1.718974029
9171272)}
Sum of tuple values of likelihood dictionary for label e is (15132.5, 1.0, -108.68107160193631)

```

The formula for the calculation of likelihood is in the image where we first compute the sum of frequencies for a character and add $\alpha = 0.5$ and then divide it by sum of the counts across all letters and space and add $27 * \alpha$.

$$\text{prob} = (\text{columnSum} + \alpha) / (\text{sum of all columns} + (\text{number of columns}) * (\alpha))$$

The likelihood vector for e looks something like this: ['a': 0.0601685114819098, 'b': 0.011134974392863043, 'c': 0.021509995043779945, 'd': 0.021972575582355856, 'e': 0.1053692383941847, 'f': 0.018932760614571286, 'g': 0.017478936064761277, 'h': 0.047216256401784236, 'i': 0.055410540227986124, 'j': 0.001420783082768875, 'k': 0.0037336857756484387, 'l': 0.028977366595076822, 'm': 0.020518751032545846, 'n': 0.057921691723112505, 'o': 0.06446390219725756, 'p': 0.01675202378985627, 'q': 0.0005617049396993227, 'r': 0.053824549810011564, 's': 0.06618205848339666, 't': 0.08012555757475633, 'u': 0.026664463902197257, 'v': 0.009284652238559392, 'w': 0.015496448042293078, 'x': 0.001156451346439782, 'y': 0.013844374690236246, 'z': 0.0006277878737815959, 'space': 0.1792499586981662]

- Print θ_j, θ_s and include in final report the class conditional probabilities for Japanese and Spanish. For Japanese, we get the following values ['a': 0.1317656102589189, 'b': 0.010866906600510151, 'c': 0.005485866033054963, 'd': 0.01722631818022992, 'e': 0.06020475907613823, 'f': 0.003878542227191726, 'g': 0.014011670568503443, 'h': 0.03176211607673224, 'i': 0.09703343932352633, 'j': 0.0023411020650616725, 'k': 0.05740941332681086, 'l': 0.001432614696530277, 'm': 0.03979873510604843, 'n': 0.05671057688947902, 'o': 0.09116321324993885, 'p': 0.0008735455466648031, 'q': 0.00010482546559977637, 'r': 0.04280373178657535, 's': 0.0421747789929767, 't': 0.056990111464411755, 'u': 0.07061742199238269, 'v': 0.0002445927530661449, 'w': 0.01974212935462455, 'x': 3.4941821866592126e-05, 'y': 0.01415143785596981, 'z': 0.00772214263251686, 'space': 0.12344945665466997] For Spanish we get the following values, ['a': 0.10456045141993771, 'b': 0.008232863618143134, 'c': 0.03752582405722919, 'd': 0.039745922111559924, 'e': 0.1138108599796491, 'f': 0.00860287996053159,

'g': 0.0071844839813758445, 'h': 0.0045327001942585795, 'i': 0.049859702136844375, 'j': 0.006629459467793161, 'k': 0.0002775122567913416, 'l': 0.052943171656748174, 'm': 0.02580863988159477, 'n': 0.054176559464709693, 'o': 0.07249236841293824, 'p': 0.02426690512164287, 'q': 0.007677839104560451, 'r': 0.05929511886774999, 's': 0.06577040485954797, 't': 0.03561407295488884, 'u': 0.03370232185254849, 'v': 0.00588942678301625, 'w': 9.250408559711388e-05, 'x': 0.0024976103111220747, 'y': 0.007862847275754679, 'z': 0.0026826184823163022, 'space': 0.16826493170115014]

4. Treat e10.txt as a test document x . Represent x as a bag-of-words count vector (Hint: the vocabulary has size 27). Print the bag-of-words vector x and include in final report.

For the bag of characters model we need to have the frequencies of each character in the file. Bag of characters vector for e10.txt is

a b c d e f g h i j k l m n o p q r s t u v w x y z space

164 32 53 57 311 55 51 140 140 3 6 85 64 139 182 53 3 141 186 225 65 31 47 4 38 2 498

5. Compute $\hat{p}(x | y)$ for $y = e, j, s$ under the multinomial model assumption, respectively. Use the formula

$$\hat{p}(x | y) = \prod_{i=1}^d \theta_{i,y}^{x_i}$$

where $x = (x_1, \dots, x_d)$. Show the three values: $\hat{p}(x | y = e), \hat{p}(x | y = j), \hat{p}(x | y = s)$. Hint: you may notice that we omitted the multinomial coefficient. This is ok for classification because it is a constant w.r.t. y .

$\log \hat{p}(x | y = e)$ = Log sum and probability value for e conditional for the file e10.txt is -7841.865447060571, 0.0.

$\log \hat{p}(x | y = j)$ = Log sum and probability value for j conditional for the file e10.txt is -8771.433079075021, 0.0.

$\log \hat{p}(x | y = s)$ = Log sum and probability value for s conditional for the file e10.txt is -8467.282044010564, 0.0.

```

main.py > predict > log_estimated_likelihood_test_point_e10_econditional
91
92
93 def get_test_document_statistics(directory: string, filename: string, df: pd.DataFrame, conditional: pd.DataFrame):
94     df = df.loc[df['file_name'] == filename]
95     df = df.iloc[:, :27]
96     # print('Bag of characters vector for', filename,
97     #       'is', df.to_string(index=False))
98     log_sum = 0
99     prob = 1
100     with open(directory+filename) as fileObj:
101         for line in fileObj:
102             for ch in line:
103                 if ch == '\n':
104                     continue
105                 elif ch == ' ':
106                     log_sum += conditional['space'][2]
107                     prob *= conditional['space'][1]
108                 else:
109                     log_sum += conditional[ch][2]
110                     prob *= conditional[ch][1]
111     print('Log sum and probability value for ', filename, 'is', log_sum, prob)
112     return log_sum, prob
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000

```

Log sum and probability value for e17.txt is -7355.874463641964 0.0
Log sum and probability value for e13.txt is -4742.753503541296 0.0
Log sum and probability value for e13.txt is -5182.133348783306 0.0
Log sum and probability value for e13.txt is -5178.663420576567 0.0
Log sum and probability value for e12.txt is -5285.490455003663 0.0
Log sum and probability value for e12.txt is -5836.362798159375 0.0
Log sum and probability value for e12.txt is -5681.093710850036 0.0
Log sum and probability value for e10.txt is -7841.865447060571 0.0
Log sum and probability value for e10.txt is -8771.433079075021 0.0
Log sum and probability value for e10.txt is -8467.282044010564 0.0
Log sum and probability value for e11.txt is -9346.211875059087 0.0
Log sum and probability value for e11.txt is -10407.2838331908 0.0
Log sum and probability value for e11.txt is -10056.252113896917 0.0

6. Use Bayes rule and your estimated prior and likelihood, compute the posterior $\hat{p}(y | x)$. Show the three values: $\hat{p}(y = e | x)$, $\hat{p}(y = j | x)$, $\hat{p}(y = s | x)$. Show the predicted class label of x .

$$\hat{p}(y = e | x) = 1/e^{7841}$$

$$\hat{p}(y = j | x) = 1/e^{8771} \quad \hat{p}(y = s | x) = 1/e^{8467}$$

```

main.py> get_test_document_statistics
get_test_document_statistics(filename: string):
133
134     # q4, q5
135     log_estimated_likelihood_test_point_e10_econditional, _ = get_test_document_statistics(
136         directory, filename, test, e_conditional)
137     log_estimated_likelihood_test_point_e10_jconditional, _ = get_test_document_statistics(
138         directory, filename, test, j_conditional)
139     log_estimated_likelihood_test_point_e10_sconditional, _ = get_test_document_statistics(
140         directory, filename, test, s_conditional)
141
142     # q6
143     log_posterior_e = log_estimated_likelihood_test_point_e10_econditional + log_prior_e
144     log_posterior_j = log_estimated_likelihood_test_point_e10_jconditional + log_prior_j
145     log_posterior_s = log_estimated_likelihood_test_point_e10_sconditional + log_prior_s
146
147     print('Log posterior values for ', filename, 'is', log_posterior_e,
148         log_posterior_j, log_posterior_s)
149     max_posterior = max(log_posterior_e, log_posterior_j, log_posterior_s)
150     pred_label = 'e' if max_posterior is log_posterior_e else 's' if max_posterior is log_posterior_s else 'j'
151     print('Predicted label for filename ', filename, 'is ', pred_label, 'having posterior log probability as ',
152         max_posterior, 'which means a probability equal to 1/e^', abs(max_posterior))
153     return pred_label
154     pred_label_e10 = predict('e10.txt')
155     print('Pred label for e10.txt is', pred_label_e10)

```

```

PROBLEMS 1 OUTPUT DEBUG CONSOLE TERMINAL PORTS
[{'a': 0.10456045141993771}, {'b': 0.008232863618143134}, {'c': 0.03752582405722}, {'d': 0.00860287996053159}, {'e': 0.0071844839813758445}, {'f': 0.004532700194}, {'g': 0.0002775122567913416}, {'h': 0.052943171656748174}, {'i': 0.02580863988159477}, {'j': 0.054176559464709693}, {'k': 0.07249236}, {'l': 0.0426690512164287}, {'m': 0.007677839104560451}, {'n': 0.05929511886774999}, {'o': 0.06577040485954797}, {'p': 0.0356140}, {'q': 0.03370232185254849}, {'r': 0.00588942678301625}, {'s': 9.250408559711388e-05}, {'t': 0.0024976103111220747}, {'u': 0.007}, {'v': 0.0026826184823163022}, {'w': 0.16826493170115014}, {'x': 0.0026826184823163022}, {'y': 0.0026826184823163022}, {'z': 0.0026826184823163022}, {'space': 0.16826493170115014}]
Log sum and probability value for e10.txt is -7841.865447060571 0.0
Log sum and probability value for e10.txt is -8771.433079075021 0.0
Log sum and probability value for e10.txt is -8467.282044010564 0.0
Log posterior values for e10.txt is -7842.964059349239 -8772.531691363689 -8468.380656299232
Predicted label for filename e10.txt is e having posterior log probability as -7842.964059349239 which means a probability equal to 1/e^
7842.964059349239

```

7. Evaluate the performance of your classifier on the test set (files 10.txt to 19.txt in three languages). Present the performance using a confusion matrix. A confusion matrix summarizes the types of errors your classifier makes, as shown in the table below. The columns are the true language a document is in, and the rows are the classified outcome of that document. The cells are the number of test documents in that situation. For example, the cell with row = English and column = Spanish contains the number of test documents that are really Spanish, but misclassified as English by your classifier.

		ACTUAL		
		e	j	s
PREDICTED	e	10	0	0
	j	0	10	0
	s	0	0	10

30 test documents from e10-e19.txt, j10-j19.txt and s10-s19.txt. This means that the accuracy of Naive Bayes classifier is a whopping 100%!

```

148     log_posterior_j, log_posterior_s)
149     max_posterior = max(log_posterior_e, log_posterior_s)
150     pred_label = 'e' if max_posterior is log_posterior_e else 's'
151     print('Predicted label for filename ', filename, ' is ', pred_label)
152     return max_posterior, 'which means a probability of ', max_posterior
153     return pred_label
154     pred_label_e10 = predict('e10.txt')
155     print('Pred label for e10.txt is', pred_label_e10)
156
157     # q7
158     labels = ['e', 'j', 's']
159     # column: actual class, row : pred class
160     confusion_matrix = pd.DataFrame(0, index=labels, columns=labels)
161     for file in test['file_name']:
162         y_actual = file[0]
163         y_pred = predict(file)
164         confusion_matrix[y_pred][y_actual] += 1
165     print(confusion_matrix)
166

```

```

Predicted label for filename s18.txt is s having posterior log
5279.89826423505
Log sum and probability value for s19.txt is -3373.41911356256
Log sum and probability value for s19.txt is -3744.13497713241
Log sum and probability value for s19.txt is -3202.78604695219
Log posterior values for s19.txt is -3374.51772585125 -3745.23
Predicted label for filename s19.txt is s having posterior log
3203.884659240867

```

	e	j	s
e	10	0	0
j	0	10	0
s	0	0	10

8. If you take a test document and arbitrarily shuffle the order of its characters so that the words (and spaces) are scrambled beyond human recognition. How does this shuffling affect your Naive Bayes classifier's prediction on this document? Explain the key mathematical step in the Naive Bayes model that justifies your answer.

I believe that it will not affect the Naive Bayes classifier's prediction on the document. This stems from the fundamental principle and assumption about Naive Bayes that each character is independent of one another due to which we multiply the probability of each terms. It's only the presence or absence of specific characters that are affecting the probability value.

The Naive Bayes classifier doesn't consider word ordering between words in the document.

The key mathematical step that justifies this is the application of the chain rule of probability. According to this rule:

$$P(L, C_1, C_2, \dots, C_n) = P(L) \times P(C_1|L) \times P(C_2|L, C_1) \times \dots \times P(C_n|L, C_1, C_2, \dots, C_{n-1})$$

Where: - L is the class label, - C_1, C_2, \dots, C_n are the chars in the document.

Naive Bayes simplifies this by assuming that each word's occurrence is independent of the others given the class label. Therefore, the above equation becomes:

$$P(L, C_1, C_2, \dots, C_n) = P(L) \times P(C_1|L) \times P(C_2|L) \times \dots \times P(C_n|L)$$

Changing the order of characters in a document will simply change the order of terms in the multiplication and will not affect the final result, hence shuffling would result in no effect.

4 Simple Feed-Forward Network (20pts)

In this exercise, you will derive, implement back-propagation for a simple neural network and compare your output with some standard library's output. Consider the following 3-layer neural network.

$$\hat{y} = f(x) = g(W_2 \sigma(W_1 x))$$

Suppose $x \in \mathbb{R}^d$, $W_1 \in \mathbb{R}^{d_1 \times d}$, and $W_2 \in \mathbb{R}^{k \times d_1}$ i.e. $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$. Let $\sigma(z) = [\sigma(z_1), \dots, \sigma(z_n)]$ for any $z \in \mathbb{R}^n$ where $\sigma(z) = \frac{1}{1+e^{-z}}$ is the sigmoid (logistic) activation function and $g(z_i) = \frac{\exp(z_i)}{\sum_{i=1}^k \exp(z_i)}$ is the softmax function. Suppose the true pair is (x, y) where $y \in \{0, 1\}^k$ with exactly one of the entries equal to 1, and you are working with the cross-entropy loss function given below,

$$L(x, y) = - \sum_{i=1}^k y \log(\hat{y})$$

1. Derive backpropagation updates for the above neural network. (5 pts)

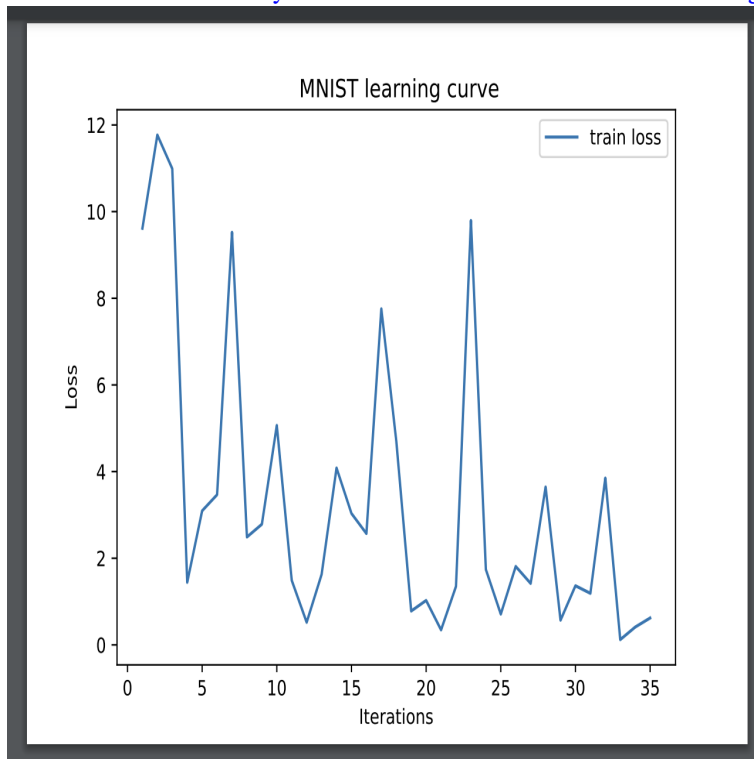
$$\frac{\partial L}{\partial W_{11}^{(2)}} = \frac{\partial L}{\partial \hat{y}_1^{(3)}} \times \frac{\partial a_1^{(3)}}{\partial z_1^{(3)}} \times \frac{\partial z_1^{(3)}}{\partial W_{11}^{(2)}}$$

where $\frac{\partial a}{\partial z} = \sigma(z) * (1 - \sigma(z))$ and $\frac{\partial z_1}{\partial W_{11}} = x_1$

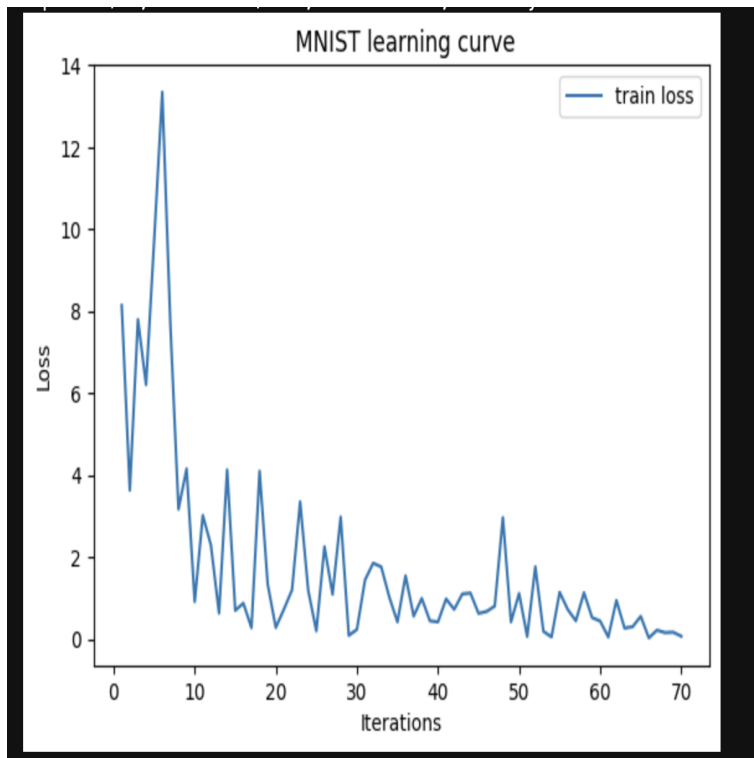
$$\frac{\partial L}{\partial W_{11}} = \frac{\partial L}{\partial a_{11}} \times \frac{\partial a_{11}}{\partial W_{11}} = (\hat{y} - y) * (W_{11}^2) * (a_{11})(1 - a_{11}) * x_1$$

Implement it in NumPy or PyTorch using basic linear algebra operations. (e.g. You are not allowed to use autograd, built-in optimizer, model, etc. in this step. You can use library functions for data loading, processing, etc.). Evaluate your implementation on MNIST dataset, report test errors and learning curve. (10 pts)

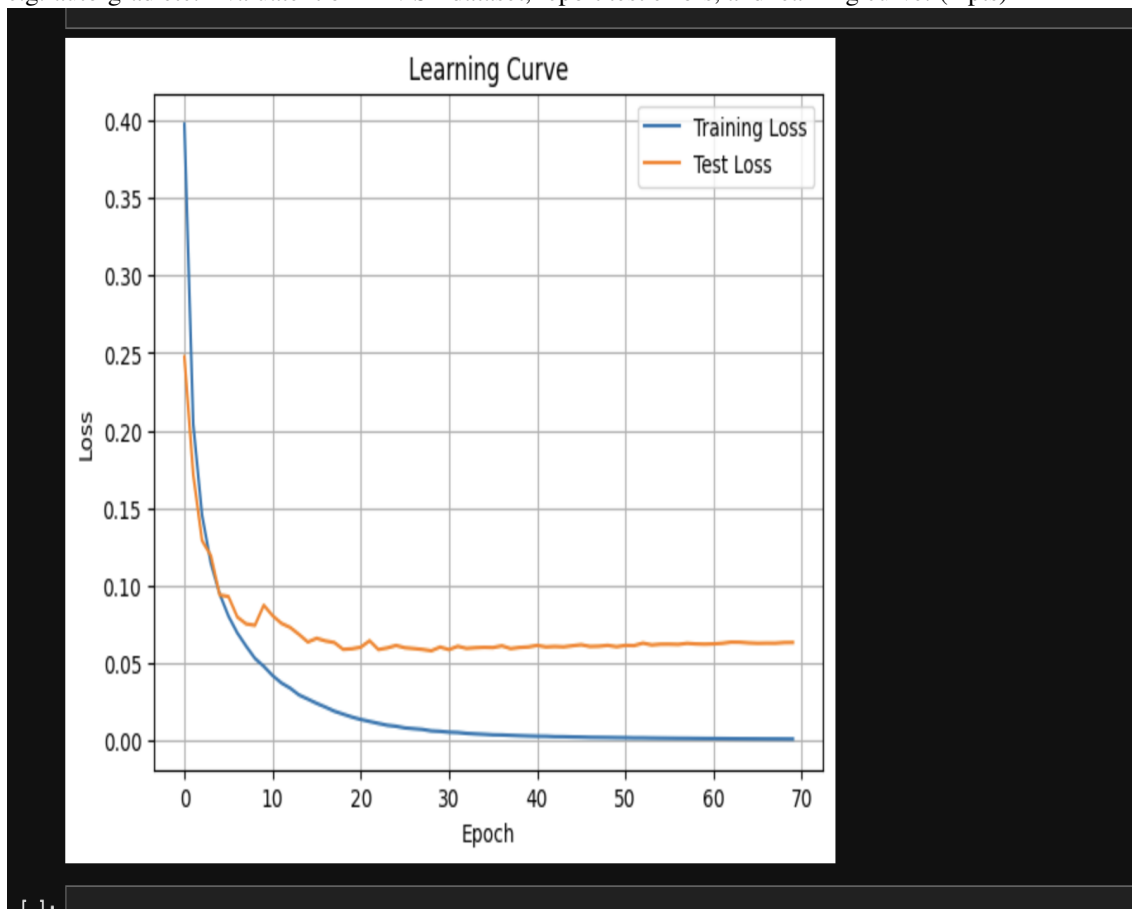
On training the model for 35 iterations with weights initialized randomly , learning rate as 0.1 and batch size of 32 we receive an accuracy of 0.83 and here is the chart for the training loss



On training the model for 70 iterations with weights initialized randomly , learning rate as 0.1 and batch size of 32 we receive an accuracy of 0.92 and here is the chart for the training loss



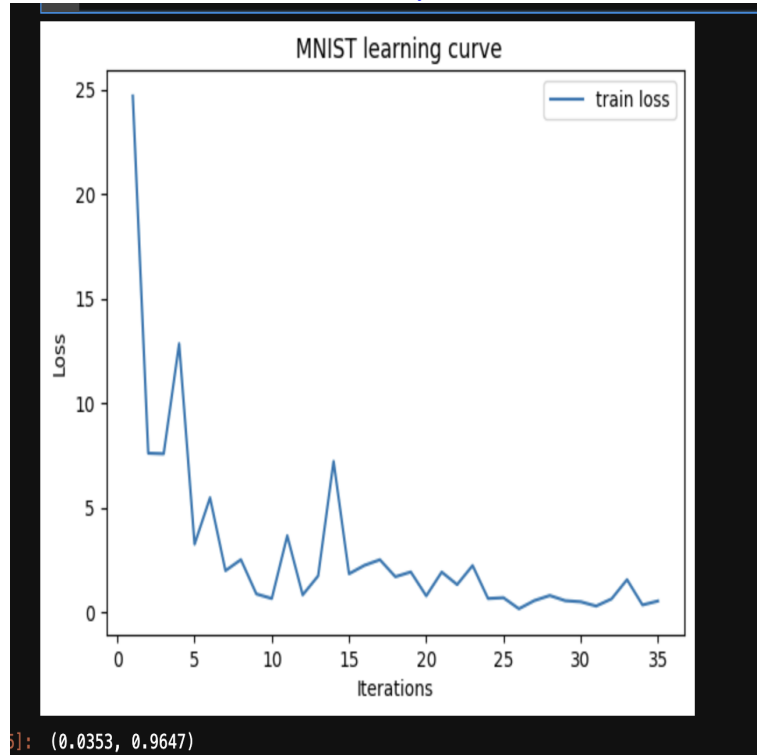
Implement the same network in PyTorch (or any other framework). You can use all the features of the framework e.g. auto-grad etc. Evaluate it on MNIST dataset, report test errors, and learning curve. (2 pts)



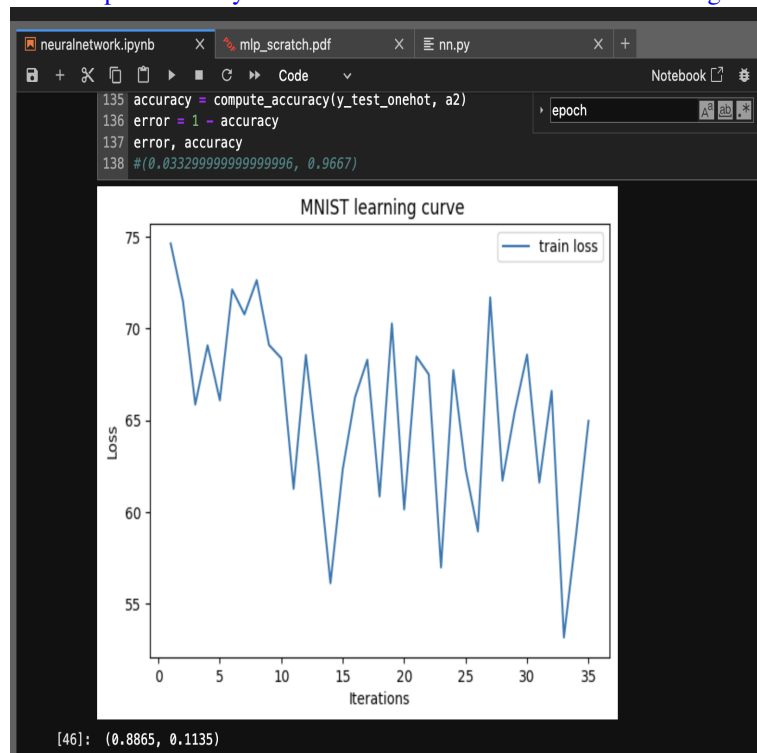
Try different weight initialization a) all weights initialized to 0, and b) initialize the weights randomly between -1 and 1. Report test error and learning curves for both. (You can use either of the implementations) (3 pts)

On training the model for 35 iterations with weights initialized randomly from -1 to 1, learning rate as 0.1 and

batch size of 32 we receive an accuracy of 0.96 and here is the chart for the training loss



On training the model for 35 iterations with weights initialized as 0, learning rate as 0.1 and batch size of 32 we see a drop in accuracy of 0.88 and here is the chart for the training loss



Note: The results have random initialization of weights in a few cases and hence cannot be compared. Also, a random permutation of the dataset is taken which might affect the batch size of 32 loss computation.

You should play with different hyperparameters like learning rate, batch size, etc. for your own learning. You only need to report results for any particular setting of hyperparameters. You should mention the values of those along with the results. Use $d_1 = 300$, $d_2 = 200$. For optimization use SGD (Stochastic gradient descent) without momentum, with some batch size say 32, 64, etc. MNIST can be obtained from here (<https://pytorch.org/vision/stable/datasets.html>)